

Maschine Learning

Algorithm learns class of tasks, measured by loss function, from experience.

supervised learning learn $h : \Delta^* \rightarrow \Sigma^*$, $h = t$; example: $(x, y) \in \Delta^* \times \Sigma^*$, $t(x) = y$.

unsupervised learning learn $h : \Delta^* \rightarrow \Sigma^*$, $\ker(h) = \ker(t)$; example: $x \in \Delta^*$.

reinforcement learning learn strategy based on feedback from environment.

2 Supervised Learning

- model function $t : \mathcal{M} \rightarrow \mathcal{R}$

- $\text{supp}(t) = \{m \in \mathcal{M} \mid t(m) \neq 0\}$

- $\bar{m} \in \text{supp}(t) \Leftrightarrow t(\bar{m}) = 1$

Hypothesis of A: potential result of A

Hypothesis space \mathcal{H}_A of A: set of all hypotheses

h fits D if $h(x_i) = y_i$ for all $(x_i, y_i) \in D$

Version space $\mathcal{V}_A(D)$ of A: all hypotheses that fit D

Inductive bias of A: set of assumptions that A uses to predict outputs of unseen data

2.1 Conjunctive Clause

$\theta = (\theta_1, \dots, \theta_k), \theta_i \in M_i \cup \{\star, \perp\}$

- $\theta_\perp = (\perp, \dots, \perp)$ most specific

- $\theta_\star = (\star, \dots, \star)$ most general

- $\text{supp}(h_{\theta_\perp}) = \emptyset, \text{supp}(h_{\theta_\star}) = \mathcal{M}$

- $h_{\theta_\perp} = h_{(\theta_1, \dots, \perp, \dots, \theta_k)}$...

induced hypothesis $h_\theta(m_1, \dots, m_k) = 1$ if $\forall i : \theta_i \in \{m_i, \star\}$ else 0

$h \preceq h'$ if $\text{supp}(h) \subseteq \text{supp}(h')$. h is more specific (less general) than h'

Find-S Algorithm finds most specific conjunctive clause that fits D

1. Start with $\theta_\perp = (\perp, \dots, \perp)$
2. iterate over POSITIVE examples
3. min-generalize θ to fit example
4. $\perp \rightarrow a, a \rightarrow \star$

- maximal general hypothesis:

1. start at $\theta_\star = (\star, \dots, \star)$
2. exclude every negative example
3. $(\star, \dots) \rightarrow \{(b, \dots), (c, \dots)\}$

- If $\mathcal{V}_A(D) \neq \emptyset$, Find-S finds $h \in \mathcal{V}_A(D)$

disjunctive normal form $\Theta = \{\theta_1, \dots, \theta_m\}$

'finite set of conjunctive clauses'

induced hypothesis $h_\Theta(\bar{m}) = 1$ if $\exists \theta \in \Theta : h_\theta(\bar{m}) = 1$ else 0

- $\text{supp}(\Theta) = \bigcup_{\theta \in \Theta} \text{supp}(\theta)$

- can represent all boolean functions

Boundary sets of version space

maximally general hypotheses $V_A^\top(D) = \{h \in V_A(D) \mid \nexists h' \in V_A(D) : h \prec h'\}$

maximally specific hypotheses $V_A^\perp(D) = \{h \in V_A(D) \mid \nexists h' \in V_A(D) : h' \preceq h\}$

- $h \in V_A^\top(D)$ maximal, weil: $\forall x \in M \setminus \text{supp}(h) : \text{supp}(h) \cup \{x\} \notin \text{supp}(V_A(D))$

Theorem: $V_A(D) = \{h \in H_A \mid \exists h_\top \in V_A^\top(D), \exists h_\perp \in V_A^\perp(D) : h_\perp \preceq h \preceq h_\top\}$

$\rightarrow V_A(D)$ det. by $V_A^\top(D)$ and $V_A^\perp(D)$

- only 1 lower bound (in $V_A^\perp(D)$), potentially multiple upper bounds (in $V_A^\top(D)$)

Candidate Elimination Algorithm

Output: DNF for $V_A^\top(D)$ and $V_A^\perp(D)$

1. $S_\perp = \{\theta_\perp\}, S_\top = \{\theta_\star\}$

2. for $1 \leq i \leq n : y_i = 1$ (pos. xmpls)

1. keep only fitting h from S_\top

2. $\forall \theta \in S_\perp : h_\theta(x_i) = 0$

- remove θ , add all min generalizations θ' of θ that fit x_i to S_\perp

3. keep only most specific h in S_\perp

3. for $1 \leq i \leq n : y_i = 0$ (neg. xmpls)

1. keep only fitting h from S_\perp

2. $\forall \theta \in S_\top : h_\theta(x_i) = 1$

- remove θ , add all min specializations θ' of θ that fit x_i to S_\top , for which a more specific $\theta_\perp \in S_\perp$ exists!

3. keep only most general h in S_\top

- $V_A^\top = \{h_\theta \mid \theta \in S_\top\}, V_A^\perp = \{h_\theta \mid \theta \in S_\perp\}$

- Concept indentified if: $S_\perp = S_\top$ and $|S_\top| = 1$. $V_A(D) = \emptyset$ if $S_\perp = \emptyset \vee S_\top = \emptyset$

2.2 Decision Trees

Splitting $\Pi = \{M_1, \dots, M_p\}$ is finite partition of (sub)feature Space \mathcal{M}'

- induces splitting of $\{1, \dots, n\}$ into $I_{D'}(M_1), \dots, I_{D'}(M_p)$ (sets of indices)

- monotonic splits: based on 1 feature

- simple split: monotonic, into all realizations $M = \{\bar{m} \in M \mid m_1 = a, \dots\}$

- binary split: monotonic, into 2 sets

$M = \{\bar{m} \in M \mid m_1 \in A\} \cup \{\bar{m} \in M \mid m_1 \notin A\}$

- induced hypothesis $h_T(\bar{m}) = T(v)$, where v is unique leaf s.t. $\bar{m} \in M_v$

- simple decision trees can represent all hypotheses

Decision Tree Quality Measures

- Number of leaves

- Height (max number of constraints to check)

- External path length (sum of all path lengths from root to leaf)

- Weighted external path length (sum of all path lengths from root to leaf, weighted by number of examples classified in that leaf)

Theorem: Given D and bound b, its NP hard to decide existence of decision tree T s.t. h_T fits D and T has ext. p.l. $\leq b$

- Majority Class $\text{Maj}_D(M')$ maj. $r \in R$

- Number of Misclassifications: $Err_D(M', r)$ in feature subspace M' with majority class r

- $Err_D(T)$: sum up all $Err_D(M_v, T(v))$

Pure Node v if $Err_D(M_v, T(v)) = 0$

- class distribution $p^{M'}_r(r) : p(r)$ in M'

- **Impurity Function** $\iota : [0, 1]^R \rightarrow \mathbb{R}$ if

- $\iota(p)$ is minimal $\forall p : p(r) = 1$
- ι symmetric in classes
- ι is maximal for uniform distr.

gets probability distribution as input

1. $\bar{p}(r) = 1 - \max_{r \in R} p(r)$

2. Entropy $H(p) = - \sum_{r \in R} p(r) \log_2 p(r)$

3. Gini Impurity $G(p) = 1 - \sum_{r \in R} p(r)^2$

- Impurity of M' is $\iota_D(M') = \iota(p^{M'}_D)$

Impurity Reduction of splitting

$\Pi = \{M'_1, \dots, M'_p\}$ of M' is:

$$\iota_D(\Pi) = \iota_D(M') - \sum_{i=1}^p \frac{|I_D(M'_i)|}{|I_D(M')|} \iota_D(M'_i)$$

Tree Construction for $M' \subseteq M$

1. if no elements in M' : new leave v : $T(v) = \text{Maj}_D(M)$

2. if $\iota(M') \leq \epsilon$: new leaf v : $T(v) = \text{Maj}_D(M')$

3. else: select split Π of M' with maximal impurity reduction

- strict imp. fct.: concave at every point

- $\iota_D(\Pi) \geq 0 \forall \Pi$ and strict imp. fct. ι

ID3 simple D.T., monothetic simple splits, impurity function: entropy.

inductive bias: local optimization (greedy)

CART D.T., binary splits, impurity function: Gini impurity

- true loss of $h \in \mathcal{H}_A$: misclassifications: $l^*(h) = \sum_{\bar{m} \in M} (1 - \delta_{h(\bar{m}), t(\bar{m})})$

- **h overfits** D if $\exists h' \in \mathcal{H}_A :$

$$l(h, D) < l(h', D) \text{ and } l^*(h) > l^*(h')$$

when: training data: noisy, small, biased

- Training Data: optimize loss here

- Validation: optimize hyperparameters

- Test Data: final estimation (true loss)

- **h overfits** (D, D_V) if $\exists h' \in \mathcal{H}_A :$

$$l(h, D) < l(h', D) \text{ and } l^*(h', D_V) < l^*(h, D_V)$$

true loss of h estimated by $l(h, D_V)$

Countermeasures to overfitting:

- increase data quality/quantity
- early stopping (no more splits)
- thrld large \rightarrow omits useful splits
- thrld small \rightarrow Large Tree
- regularization (penalize model complexity in training process)

Pruning: turn inner node v into leave with label $\text{Maj}(M_v)$

D.T. pruning Algorithm

1. given fully trained D.T.
2. prune every inner node als long as pruning doesnt increase validation loss: $l(h'_T, D_V) \leq l(h_T, D_V)$

2.3 Linear Regression

$t : M \rightarrow R, M \subseteq \mathbb{R}^k, R \subseteq \mathbb{R}$

- find $w = (w_0, \dots, w_k) \in \mathbb{R}^{k+1}$ s.t.:

$$h_w(x_1, \dots, x_k) = w_0 x_0 + \sum_{i=1}^k w_i x_i$$

approximates $t \Leftrightarrow w$ minimizes $l(h_w, D)$

- note: $x_0 = 1$ always! (w_0 is bias)

Analytical Solution

1. Partial derivatives: $\frac{\partial l(h_w, D)}{\partial w_i}$

2. Set to 0, put in values from D

3. Solve LGS with Gauss for w_i

SGD (Iterative Solution)

1. initialize w randomly

2. choose random $1 \leq i \leq n, T++$

3. $\delta = y_i - h_w(x_i)$ (residual)

4. $\Delta w = \delta \cdot x_i$ (derivatives)

5. $w = w + \eta \cdot \Delta w$ (parameters)

6. If \neg converged $\rightarrow 2$, else return w

- Pros: simple, robust to noisy data, representation independent

- Cons: stability, convergence problems, sensitive to learning rate η

BGD (accumulate derivatives $\forall i$)

1. initialize w randomly

2. For each $1 \leq i \leq n, T++$:

- 2.1. $\delta = y_i - h_w(x_i)$

- 2.2. $\Delta w = \Delta w + \delta \cdot x_i$

3. $w = w + \eta \cdot \Delta w$

4. If \neg converged $\rightarrow 2$, else return w

- sequence of examples (batch) are processed together, before updating w

IGD

1. initialize w randomly

2. For each $1 \leq i \leq n, T++$:

- 2.1. $\delta = y_i - h_w(x_i)$

- 2.2. $\Delta w = \delta \cdot x_i$

- 2.3. $w = w + \eta \cdot \Delta w$

3. If \neg converged $\rightarrow 2$, else return w

| Property | SGD stochastic | IGD iterative | BGD batch | MBGD mini-batch |
|---------------------|----------------|---------------|------------|-----------------|
| Batch size | 1 | 1 | n | varies |
| Batch selection | random | sequential | sequential | sequential |
| Parallelization | difficult | difficult | trivial | trivial |
| Space requirement | low | low | high | varies |
| Stuck local minimum | no | no | yes | varies |
| Convergence speed | slow | slow | fast | varies |

Polynomial Regression

Approach: 1. prepare nonlinear combinations of features as features (curse of dimensionality: max k features $5k \leq n$)

2. then perform linear regression on expanded feature space with SGD, BGD, IGD or analytical approach.

3. Project solution back to input (feature) space

- keep original features

- for m_i^3 also include m_i^2

- increase complexity \rightarrow increase risk of overfitting

Regularization

'penalize model complexity in training process', optimize for $l'(w, D)$:

$$l'(w, D) = l(h_w, D) + \frac{\lambda}{k} \cdot r(w)$$

- Lasso Regression: $r(w) = \sum_{i=1}^k |w_i|$

- Ridge Regression: $r(w) = \sum_{i=1}^k w_i^2$

- λ big \rightarrow more regularization \rightarrow less complex model

- k : num of features (excluding bias w_0)

Develop (S,B,I)GD for **specific loss function** $l(x)$:

1. get 1st derivative $l'(x)$ of loss:

for 1 Data example: $n = 1$, leave out \sum

2. find $-\delta = h_w(x_i) - y_i$ in 1st derivative

3. replace line 4 (derivatives) with:

$$\Delta w = l'(x) \text{ but substitute } \delta (! - 1!)$$

Logistic Regression (Classification)

'find optimal hyperplane w' by optimization, to get h_w , to get classifier:

$$h_w^c(z) = 1 \text{ if } h_w(z) \geq \frac{1}{2}, 0 \text{ otherwise}$$

- 'discriminative' classifier

- only reasonable for binary $R = \{0, 1\}$ ($|R| > 2$ induces order bias)

- Training: optimizes w for $l(h_w, D)$

- Prediction: performed by h_w^c (different)

- Discriminating Hyperplane 1 Dimension: less than w (Plane \rightarrow Line \rightarrow Point)

- logistic function: $h_w^\sigma(z) = \sigma(h_w(z))$

- logistic classifier: $h_w^{c,\sigma}(z) = 1 \text{ if } h_w^\sigma(z) \geq \frac{1}{2}, 0 \text{ otherwise}$

> MLE Maximum Likelihood Estimator given H_A for D is: $\hat{h} = \arg \max_{h \in H_A} P[D; h]$

2.4 Support Vector Machines

'learn optimal discriminating Hyperplane with maximal margin directly'
 $H(w) = \{(z_1, \dots, z_k) \in \mathbb{R}^k | w_0 + w_1 x_1 + \dots + w_k x_k = 0\}$

- Normal Representation $w = (w_0, \dots, w_k)$ is normal to discriminating Hyperplane $H(w)$

- H_1 closest x_i with $y_i = 1$ ($wz^T = 1$)

- H_0 closest x_i with $y_i = 0$ ($wz^T = -1$)

- Hyperplanes H_1, H_0 parallel to $H(w)$

Margin distance(H_1, H_0) = $\frac{2}{\|w\|}$

- Hinge Loss: only falsely classified data causes loss

Hard Margin SVM: no misclassifications/boundary violations

- $\hat{w} = \arg \min_w \frac{1}{2} \vec{w} \vec{w}^T, l_h(h_w, D) = 0$

Soft Margin SVM: λ trades margin size against boundary violations

λ small: larger margin, more violations

λ big: smaller margin, less violations

- $\hat{w} = \arg \min_w \frac{1}{2} \vec{w} \vec{w}^T + \lambda l_h(h_w, D)$

Kernel Trick: Kernels permits nonlinear separation in input space \mathbb{R}^k (through linear separation in \mathbb{R}^{k+d})

- with suitable Kernel: no actual computations in \mathbb{R}^{k+d}

- > no additional effort for non linear classification (linear classifier for free)

D Linearly Separable if:

$\exists w_0, w_1, \dots, w_k : y'_i \cdot (w_0 + w_1 x_1 + \dots + w_k x_k) > 0$

where: $y'_i = 1$ if $(y_i = 1)$, -1 if $(y_i = 0)$

1. $\forall i$ where $y_i = 1 : h(x_i) > 0$

2. $\forall i$ where $y_i = 0 : h(x_i) < 0$

Proof by finding w , then transform to discriminating hyperplane (eg point x):
 $- w_0 + w_1 x_1 = 0 \Rightarrow x = -\frac{w_0}{w_1}$

Disprove by finding contradiction in System of inequations.

loss functions (and derivatives)

- $l_{0/1}(h, D) = \sum_{i=1}^n (1 - \delta_{y_i, h(x_i)})$
 $\rightarrow \delta_{ij} = 1$ if $i = j, 0$ otherwise.
- $l_2(h, D) = \frac{1}{2n} \sum_{i=1}^n (h(x_i) - y_i)^2$ (Mean Squared Error)
- $\frac{\partial l(h, D)}{\partial w_p} = \frac{1}{n} \sum_{i=1}^n (h(x_i) - y_i) x_p = \frac{1}{n} \sum_{i=1}^n (w_0 + w_1 x_{i1} + \dots + w_p x_{ip} - y_i) x_p$
- $l'(w, D) = \frac{1}{2n} \sum_{i=1}^n (h(x_i) - y_i)^2 + \frac{\lambda}{k} \sum_{i=1}^k w_i^2$ (Ridge Regression)

2.5 Neural Networks

Perception Hypothesis

$$h_w^H(z) = H(\sum_{i=0}^k w_i z_i) = 1 \text{ if } wz^T \geq 0$$

Perceptron Training

- initialize w randomly, $T = 0$
- Select random $1 \leq i \leq n$
- $\delta = y_i - h_w^H(x_i)$
- $\Delta w = \delta \cdot x_i$
- $w = w + \eta \cdot \Delta w$
- If $l(h_w^H, D) \neq 0$ to 2, else return w

Rosenblatt: If D linearly separable: PT terminates after finitely many corrections

| Property | Gradient descent | PT algorithm |
|--------------------------------------|--------------------------------------|----------------------------------|
| Loss function Discriminator | ℓ_2 or ℓ_∞ hyperplane | 0-1 loss $\ell_{0/1}$ hyperplane |
| Data inseparable Perfect separation | robust potentially | no termination guaranteed |
| Parameter updates on Correction size | surrogate error scaled by steepness | class error fixed |

Nonlinear separation possible through network of perceptrons

| Activation | Function | Symbol | Algorithm |
|------------|--|------------|---------------------|
| Linear | $id(z) = z$ | Σf | Linear regression |
| ReLU | $ReLU(z) = \max(0, z)$ | Σf | - |
| Leaky ReLU | $r_a(z) = \begin{cases} \max(0, z) \\ + \alpha \cdot \min(0, z) \end{cases}$ | Σf | - |
| Heaviside | $H(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{otherwise} \end{cases}$ | Σf | PT algorithm |
| Sigmoid | $\sigma(z) = \frac{1}{1+e^{-z}}$ | Σf | Logistic regression |

Network $N = (V, E, wt)$, $wt : E \rightarrow \mathbb{R}$

- State: $s : V \rightarrow \mathbb{R}$

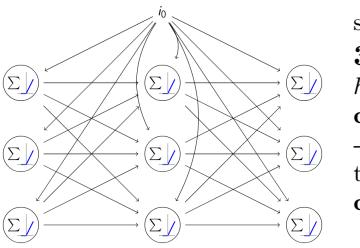
$$s'(v') = a(\sum_{v \in V} s(v) \cdot wt(v, v'))$$

a (weighted sum of previous states)

- OR: $N = W \in \mathbb{R}^{(V \cup I) \times V}$ (adj matrix)

- FeedForward N: no loops, no cycles, every node reachable from 1 input node

Example



$w_0 * i_0 = w_0$ is bias, noted above node

Characteristics: 1. input, 2. internal nodes, 3. layers, 4. type (FFN/RNN?)

Induced Regression Hypothesis

$$h_w^v(z) : \mathbb{R}' \rightarrow \mathbb{R}, h_w^v(z) = s_v^{(m)}$$

Forward Pass with fixed output node

Neural Network Training

compute gradients gradients

> Gradients of Nodes $\frac{\partial L}{\partial r_v} =$

- $(s_v - y) \cdot H(r_v)$ if v designated o node
- 0 if v other o node
- $H(r_v) \cdot \sum_{v' \in V} W_{vv'} \cdot \frac{\partial L}{\partial r_{v'}}$ otherwise

> Gradients of Edges $\frac{\partial L}{\partial W_{vv'}}$: 'State of predecessor · gradient of successor'

FF NN Training with SGD

- initialize W randomly
- choose random $1 \leq i \leq n, T++$
- compute states s_v (forward pass)
- compute gradients ΔL for (x_i, y_i) (backward pass)
- $W = W - \eta \cdot \Delta L$ (update weights)
- If -converged $\rightarrow 2$, else return W

- RNNs good for sequential data, gives 'sequence regression hypothesis'

m-unroll of RNN: FNN with m layers. backpropagation through time: unroll RNN, compute gradients for each layer, then average gradients for each parameter over all layers

Transformer

'multiclass classification with attention'

- softmax $\sigma : \mathbb{R}^V \rightarrow \mathbb{R}^V$ outputs distribution over inputs

- input \rightarrow (word) embedding \rightarrow encoder \rightarrow attention layer \rightarrow decoder \rightarrow output

- FF attention layer indicates relevant source parts

3 Unsupervised Learning

$h : M \rightarrow \{1, \dots, c\}$ assigns each point a **cluster**: names $(1, \dots, c)$ arbitrary

- loss is avg of squared euclidian distance to cluster mean (centroid)

c-means (Lloyd's) Algorithm

- init c centroids m_1, \dots, m_c random
- Add each i of x_i to closest centroids (m_j) Set I_j
- Update centroids to mean of assigned points. $m_j = \frac{1}{n} \sum_{i \in I_j} x_i$

- induced hypothesis: $h_{m_1, \dots, m_c}(x) = \arg \min_{1 \leq j \leq c} \|x - m_j\|^2$

- c needs to be known in advance: elbow method: plot loss for $1 \leq c \leq c_{\max}$, look for 'elbow' where loss reduction flattens

DBSCAN (Density Based Spatial Clustering of Applications with Noise)

- Hyperparameters: ϵ (radius), c_{\min} (min points in ϵ -neighborhood to be core)

DBSCAN algorithm

[Ester, Kriegel, Sander, Xu 1996]

- Compute $C \leftarrow \{x_i \mid 1 \leq i \leq n, x_i \text{ is core}\}$ (initialize core points)
- $\alpha \leftarrow 1$ and $B \leftarrow C$ (first cluster and need to handle all core points)
- While $B \neq \emptyset$
 - Select $x \in B$, $C_o \leftarrow \emptyset$ and $C_o \leftarrow \{x\}$ (select core point and setup cluster)
 - While $C_o \neq C_o'$ (as long as cluster grows)
 - $C_o \leftarrow C_o \cup C_o'$ (expand cluster)
 - $C_o' \leftarrow C_o \mid \{x_i \mid 1 \leq i \leq n, \exists x' \in C_o: \|x_i - x'\| \leq \epsilon\}$ (determine neighbors)
 - $B \leftarrow B \setminus C_o$ and $\alpha \leftarrow \alpha + 1$ (cluster handled and prepare next cluster)
 - Output clusters $C_1, \dots, C_{\alpha-1}$ (remaining points are outliers)

given V^* : optimal policy selects action with highest expected reward:

Policy iteration

- Q-Value with deterministic policy P : $Q_V^P(q) = \sum_{q' \in Q} p(q'|q, P(q)) \cdot (R(q, P(q), q') + \gamma V(q'))$

1. for every state compute $Q_V^P(q)$ until convergence

2. update policy to take optimal action

3. If -converged $\rightarrow 1$

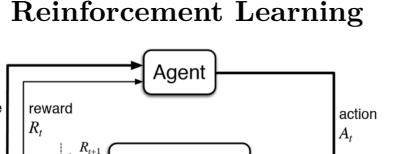
- both methods are model based: they have perfect information (exploitation)

No perfect information: model free

- ϵ -greedy policy: with prop ϵ select random action (exploration), else optimal action (exploitation)

- Q-Learning: extension of Val Iteration

- SARSA: extension of Policy Iteration



$$\begin{aligned} \frac{\partial l'(w, D)}{\partial w_p} &= \frac{1}{n} \sum_{i=1}^n (h(x_i) - y_i) x_p + 2\lambda w_p, \quad w_p \in \{0, w_1, \dots, w_k\} \quad \text{WARUM NICHT: ... + } \\ &\frac{2\lambda}{k} w_p ?!?!?!!?!?!?! \\ \bullet \quad \ell_\sigma(h, D) &= -\frac{1}{n} \sum_{i=1}^n (y_i \cdot \log(h(x_i)) + (1 - y_i) \cdot \log(1 - h(x_i))) \quad (\text{Logistic Loss}) \\ \frac{\partial \ell_\sigma(h, D)}{\partial w_p} &= -\frac{1}{n} \sum_{i=1}^n (y_i - h^\sigma(x_i)) \cdot x_ip \\ \bullet \quad l_h(h, D) &= \frac{1}{n} \sum_{i=1}^n \max(0, 1 - (2y_i - 1)h(x_i)) \quad (\text{Hinge Loss}) \\ \bullet \quad \frac{1}{n} \sum_{i=1}^n H(1_{y_i}, h(x_i)) & \quad (\text{Cross Entropy Loss}) \end{aligned}$$

$$- H(p, p') = - \sum_{r \in R} p(r) \log_2 p'(r)$$

$$\frac{\partial H(1_y, p)}{\partial s_v} = \frac{p(v)-1_y(v)}{\ln 2}$$

- $l_V(h, D) = \frac{1}{n} \sum_{i=1}^n \|x_i - \mu_i\|^2$ where μ_i is centroid of cluster that x_i belongs to (euclidean distance to cluster mean)

Ableitungsregeln

- Produktregel: $(f \cdot g)(x)' = f'(x) \cdot g(x) + f(x) \cdot g'(x)$
- Quotientenregel: $\left(\frac{f}{g}\right)'(x) = \frac{f'(x)g(x) - f(x)g'(x)}{g^2(x)}$

- Kettenregel: $(f \circ g)'(x) = (f' \circ g)(x) \cdot g'(x) = f'(g(x)) \cdot g'(x)$
 $(f \circ g \circ h)'(x) = f'(g(h(x))) \cdot g'(h(x)) \cdot h'(x)$
- $\left(\frac{1}{g}\right)'(x) = -\frac{g'(x)}{g^2(x)}$
- $\left(\frac{1}{x^n}\right)' = -nx^{-n-1}$
- $\log_a'(x) = \frac{1}{x \ln(a)}$
- $\ln_e'(x) = \frac{1}{x}$
- $|x|' = \frac{x}{|x|} = \text{sgn}(x)$ for $x \neq 0$